

Import Dataset

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
df= pd.read_csv('Heart_Attack_prediction_Dataset.csv')
df.head()
```

Out[1]:

	Patient ID	Age	Sex	Cholesterol	Blood Pressure	Heart Rate	Diabetes	Family History	Smoking	Obesity
0	BMW7812	67	Male	208	158/88	72	0	0	1	
1	CZE1114	21	Male	389	165/93	98	1	1	1	
2	BNI9906	21	Female	324	174/99	72	1	0	0	
3	JLN3497	84	Male	383	163/100	73	1	1	1	
4	GFO8847	66	Male	318	91/88	93	1	1	1	

5 rows × 26 columns



In [2]: df.columns

```
Out[2]: Index(['Patient ID', 'Age', 'Sex', 'Cholesterol', 'Blood Pressure',
       'Heart Rate', 'Diabetes', 'Family History', 'Smoking', 'Obesity',
       'Alcohol Consumption', 'Exercise Hours Per Week', 'Diet',
       'Previous Heart Problems', 'Medication Use', 'Stress Level',
       'Sedentary Hours Per Day', 'Income', 'BMI', 'Triglycerides',
       'Physical Activity Days Per Week', 'Sleep Hours Per Day', 'Country',
       'Continent', 'Hemisphere', 'Heart Attack Risk'],
      dtype='object')
```

```
In [3]: df.drop_duplicates(inplace=True)
df.shape
```

Out[3]: (8763, 26)

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8763 entries, 0 to 8762
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Patient ID      8763 non-null   object  
 1   Age              8763 non-null   int64  
 2   Sex              8763 non-null   object  
 3   Cholesterol     8763 non-null   int64  
 4   Blood Pressure   8763 non-null   object  
 5   Heart Rate       8763 non-null   int64  
 6   Diabetes         8763 non-null   int64  
 7   Family History   8763 non-null   int64  
 8   Smoking          8763 non-null   int64  
 9   Obesity          8763 non-null   int64  
 10  Alcohol Consumption 8763 non-null   int64  
 11  Exercise Hours Per Week 8763 non-null   float64 
 12  Diet              8763 non-null   object  
 13  Previous Heart Problems 8763 non-null   int64  
 14  Medication Use   8763 non-null   int64  
 15  Stress Level     8763 non-null   int64  
 16  Sedentary Hours Per Day 8763 non-null   float64 
 17  Income            8763 non-null   int64  
 18  BMI               8763 non-null   float64 
 19  Triglycerides    8763 non-null   int64  
 20  Physical Activity Days Per Week 8763 non-null   int64  
 21  Sleep Hours Per Day 8763 non-null   int64  
 22  Country           8763 non-null   object  
 23  Continent         8763 non-null   object  
 24  Hemisphere        8763 non-null   object  
 25  Heart Attack Risk 8763 non-null   int64  
dtypes: float64(3), int64(16), object(7)
memory usage: 1.7+ MB
```

In [5]: `df.describe().T`

Out[5]:		count	mean	std	min	25%	5
	Age	8763.0	53.707977	21.249509	18.000000	35.000000	54.000000
	Cholesterol	8763.0	259.877211	80.863276	120.000000	192.000000	259.000000
	Heart Rate	8763.0	75.021682	20.550948	40.000000	57.000000	75.000000
	Diabetes	8763.0	0.652288	0.476271	0.000000	0.000000	1.000000
	Family History	8763.0	0.492982	0.499979	0.000000	0.000000	0.000000
	Smoking	8763.0	0.896839	0.304186	0.000000	1.000000	1.000000
	Obesity	8763.0	0.501426	0.500026	0.000000	0.000000	1.000000
	Alcohol Consumption	8763.0	0.598083	0.490313	0.000000	0.000000	1.000000
	Exercise Hours Per Week	8763.0	10.014284	5.783745	0.002442	4.981579	10.069000
	Previous Heart Problems	8763.0	0.495835	0.500011	0.000000	0.000000	0.000000
	Medication Use	8763.0	0.498345	0.500026	0.000000	0.000000	0.000000
	Stress Level	8763.0	5.469702	2.859622	1.000000	3.000000	5.000000
	Sedentary Hours Per Day	8763.0	5.993690	3.466359	0.001263	2.998794	5.933000
	Income	8763.0	158263.181901	80575.190806	20062.000000	88310.000000	157866.000000
	BMI	8763.0	28.891446	6.319181	18.002337	23.422985	28.768000
	Triglycerides	8763.0	417.677051	223.748137	30.000000	225.500000	417.000000
	Physical Activity Days Per Week	8763.0	3.489672	2.282687	0.000000	2.000000	3.000000
	Sleep Hours Per Day	8763.0	7.023508	1.988473	4.000000	5.000000	7.000000
	Heart Attack Risk	8763.0	0.358211	0.479502	0.000000	0.000000	0.000000



Data Cleaning

```
In [6]: columns_to_drop = ['Hemisphere', 'Patient ID', 'Continent']
df.drop(columns_to_drop, axis=1, inplace=True)
df.head()
```

Out[6]:

	Age	Sex	Cholesterol	Blood Pressure	Heart Rate	Diabetes	Family History	Smoking	Obesity	Consu
0	67	Male	208	158/88	72	0	0	1	0	
1	21	Male	389	165/93	98	1	1	1	1	
2	21	Female	324	174/99	72	1	0	0	0	
3	84	Male	383	163/100	73	1	1	1	0	
4	66	Male	318	91/88	93	1	1	1	1	

5 rows × 23 columns



```
# Split 'Blood Pressure' into systolic/diastolic
if "Blood Pressure" in df.columns:
    bp_split = df["Blood Pressure"].str.split("/", expand=True)
    df["Systolic_BP"] = pd.to_numeric(bp_split[0], errors="coerce")
    df["Diastolic_BP"] = pd.to_numeric(bp_split[1], errors="coerce")
    df.drop("Blood Pressure", axis=1, inplace=True)

# Confirm final structure
print("Cleaning, rounding & normalization complete!")
print(f"Final shape: {df.shape}")
print("Data preview:")
print(df.head())
```

Cleaning, rounding & normalization complete!

Final shape: (8763, 24)

Data preview:

	Age	Sex	Cholesterol	Heart Rate	Diabetes	Family History	Smoking	\
0	67	Male	208	72	0	0	1	
1	21	Male	389	98	1	1	1	
2	21	Female	324	72	1	0	0	
3	84	Male	383	73	1	1	1	
4	66	Male	318	93	1	1	1	

	Obesity	Alcohol Consumption	Exercise Hours Per Week	...	\
0	0	0	4.168189	...	
1	1	1	1.813242	...	
2	0	0	2.078353	...	
3	0	1	9.828130	...	
4	1	0	5.804299	...	

	Sedentary Hours Per Day	Income	BMI	Triglycerides	\
0	6.615001	261404	31.251233	286	
1	4.963459	285768	27.194973	235	
2	9.463426	235282	28.176571	587	
3	7.648981	125640	36.464704	378	
4	1.514821	160555	21.809144	231	

	Physical Activity Days Per Week	Sleep Hours Per Day	Country	\
0	0	6	Argentina	
1	1	7	Canada	
2	4	4	France	
3	3	4	Canada	
4	1	5	Thailand	

	Heart Attack Risk	Systolic_BP	Diastolic_BP	
0	0	158	88	
1	0	165	93	
2	0	174	99	
3	0	163	100	
4	0	91	88	

[5 rows x 24 columns]

```
In [8]: #Check for missing values
missing_summary = df.isnull().sum()
total_missing = missing_summary.sum()

print("\nMissing values per column:")
print(missing_summary)

print("\nData types summary after cleaning:")
print(df.info())
```

Missing values per column:

Age	0
Sex	0
Cholesterol	0
Heart Rate	0
Diabetes	0
Family History	0
Smoking	0
Obesity	0
Alcohol Consumption	0
Exercise Hours Per Week	0
Diet	0
Previous Heart Problems	0
Medication Use	0
Stress Level	0
Sedentary Hours Per Day	0
Income	0
BMI	0
Triglycerides	0
Physical Activity Days Per Week	0
Sleep Hours Per Day	0
Country	0
Heart Attack Risk	0
Systolic_BP	0
Diastolic_BP	0

dtype: int64

Data types summary after cleaning:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8763 entries, 0 to 8762
Data columns (total 24 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	8763 non-null	int64
1	Sex	8763 non-null	object
2	Cholesterol	8763 non-null	int64
3	Heart Rate	8763 non-null	int64
4	Diabetes	8763 non-null	int64
5	Family History	8763 non-null	int64
6	Smoking	8763 non-null	int64
7	Obesity	8763 non-null	int64
8	Alcohol Consumption	8763 non-null	int64
9	Exercise Hours Per Week	8763 non-null	float64
10	Diet	8763 non-null	object
11	Previous Heart Problems	8763 non-null	int64
12	Medication Use	8763 non-null	int64
13	Stress Level	8763 non-null	int64
14	Sedentary Hours Per Day	8763 non-null	float64
15	Income	8763 non-null	int64
16	BMI	8763 non-null	float64
17	Triglycerides	8763 non-null	int64
18	Physical Activity Days Per Week	8763 non-null	int64
19	Sleep Hours Per Day	8763 non-null	int64
20	Country	8763 non-null	object
21	Heart Attack Risk	8763 non-null	int64
22	Systolic_BP	8763 non-null	int64

```
23 Diastolic_BP           8763 non-null    int64
dtypes: float64(3), int64(18), object(3)
memory usage: 1.6+ MB
None
```

Feature Engineering

```
In [9]: # Feature Engineering
df["BMI_Stress"] = df["BMI"] * df["Stress Level"]
df["Activity_Ratio"] = df["Exercise Hours Per Week"] / (df["Sedentary Hours Per Day"]
df["BP_Product"] = df["Systolic_BP"] * df["Diastolic_BP"]
df["Sleep_Stress_Interaction"] = df["Sleep Hours Per Day"] * df["Stress Level"]
df["Substance_Use"] = df["Smoking"] * df["Alcohol Consumption"]
```

```
In [10]: print("\nData types summary after cleaning:")
print(df.info())
```

```
Data types summary after cleaning:  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 8763 entries, 0 to 8762  
Data columns (total 29 columns):  
 #   Column           Non-Null Count Dtype  
 ---  -----  
 0   Age              8763 non-null   int64  
 1   Sex              8763 non-null   object  
 2   Cholesterol      8763 non-null   int64  
 3   Heart Rate       8763 non-null   int64  
 4   Diabetes         8763 non-null   int64  
 5   Family History   8763 non-null   int64  
 6   Smoking          8763 non-null   int64  
 7   Obesity          8763 non-null   int64  
 8   Alcohol Consumption 8763 non-null   int64  
 9   Exercise Hours Per Week 8763 non-null   float64  
 10  Diet              8763 non-null   object  
 11  Previous Heart Problems 8763 non-null   int64  
 12  Medication Use   8763 non-null   int64  
 13  Stress Level    8763 non-null   int64  
 14  Sedentary Hours Per Day 8763 non-null   float64  
 15  Income            8763 non-null   int64  
 16  BMI               8763 non-null   float64  
 17  Triglycerides    8763 non-null   int64  
 18  Physical Activity Days Per Week 8763 non-null   int64  
 19  Sleep Hours Per Day 8763 non-null   int64  
 20  Country           8763 non-null   object  
 21  Heart Attack Risk 8763 non-null   int64  
 22  Systolic_BP       8763 non-null   int64  
 23  Diastolic_BP     8763 non-null   int64  
 24  BMI_Stress       8763 non-null   float64  
 25  Activity_Ratio   8763 non-null   float64  
 26  BP_Product        8763 non-null   int64  
 27  Sleep_Stress_Interaction 8763 non-null   int64  
 28  Substance_Use    8763 non-null   int64  
dtypes: float64(5), int64(21), object(3)  
memory usage: 1.9+ MB  
None
```

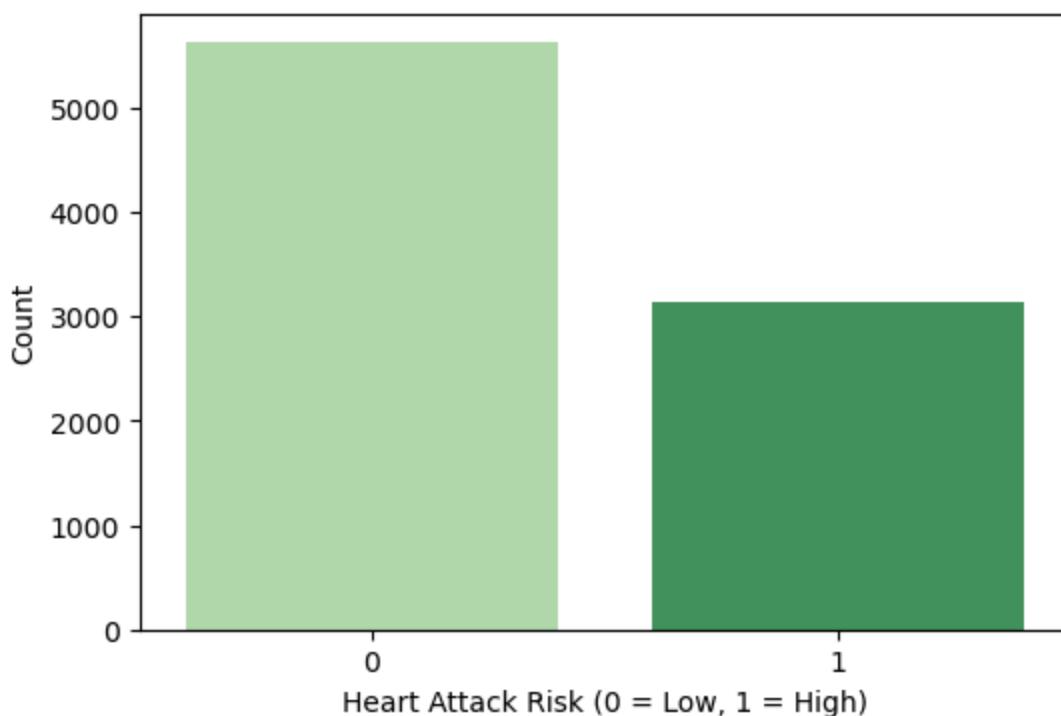
```
In [11]: df.nunique()
```

```
Out[11]: Age           73
          Sex            2
          Cholesterol    281
          Heart Rate      71
          Diabetes        2
          Family History   2
          Smoking          2
          Obesity          2
          Alcohol Consumption 2
          Exercise Hours Per Week 8763
          Diet              3
          Previous Heart Problems 2
          Medication Use     2
          Stress Level       10
          Sedentary Hours Per Day 8763
          Income             8615
          BMI                8763
          Triglycerides      771
          Physical Activity Days Per Week 8
          Sleep Hours Per Day    7
          Country            20
          Heart Attack Risk    2
          Systolic_BP         91
          Diastolic_BP        51
          BMI_Stress          8763
          Activity_Ratio       8763
          BP_Product          2835
          Sleep_Stress_Interaction 39
          Substance_Use        2
          dtype: int64
```

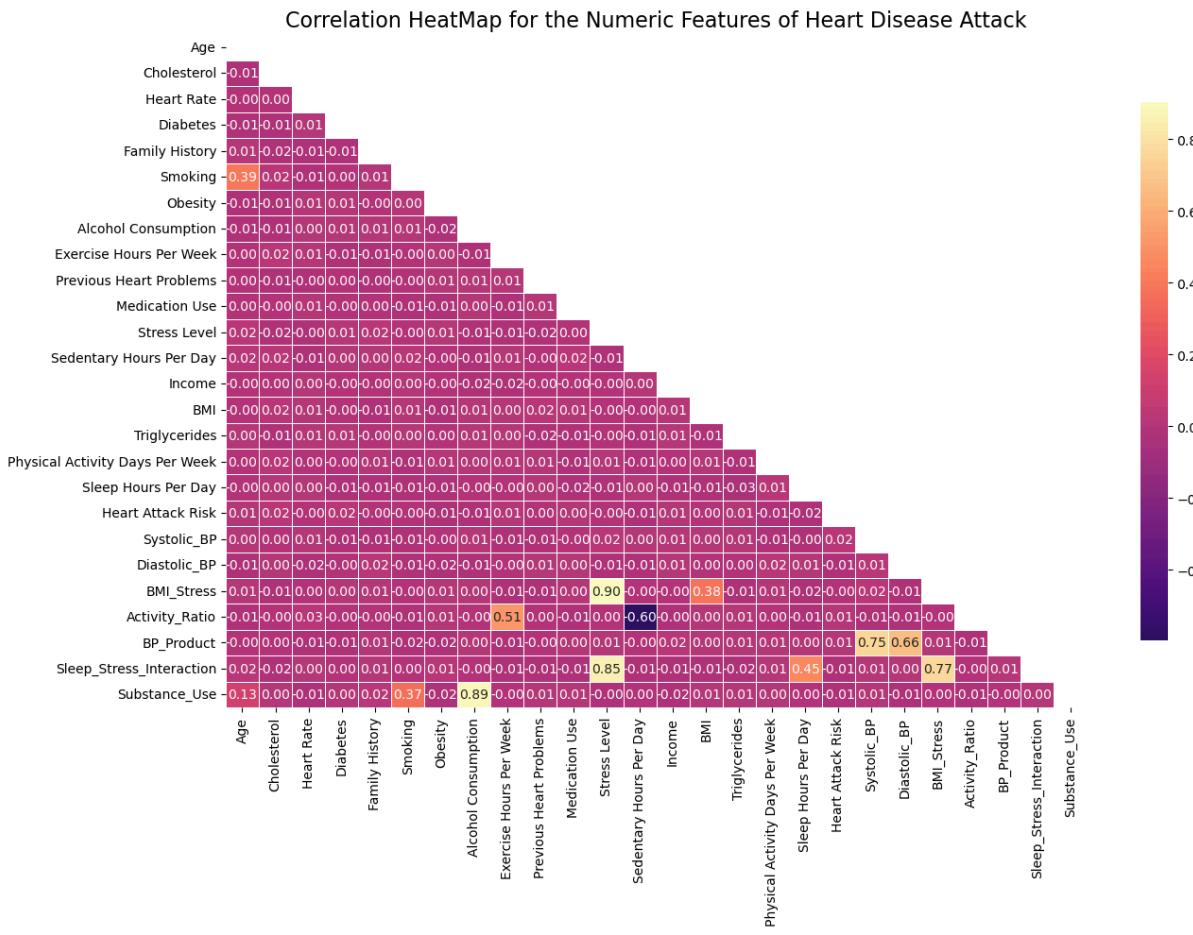
EDA and Visualizings

```
In [12]: # --- Distribution of Target Variable --
plt.figure(figsize=(6,4))
sns.countplot(x="Heart Attack Risk", hue="Heart Attack Risk", data=df, palette="Greens")
plt.title("Distribution of Heart Attack Risk", fontsize = 16)
plt.xlabel("Heart Attack Risk (0 = Low, 1 = High)")
plt.ylabel("Count")
plt.show()
```

Distribution of Heart Attack Risk

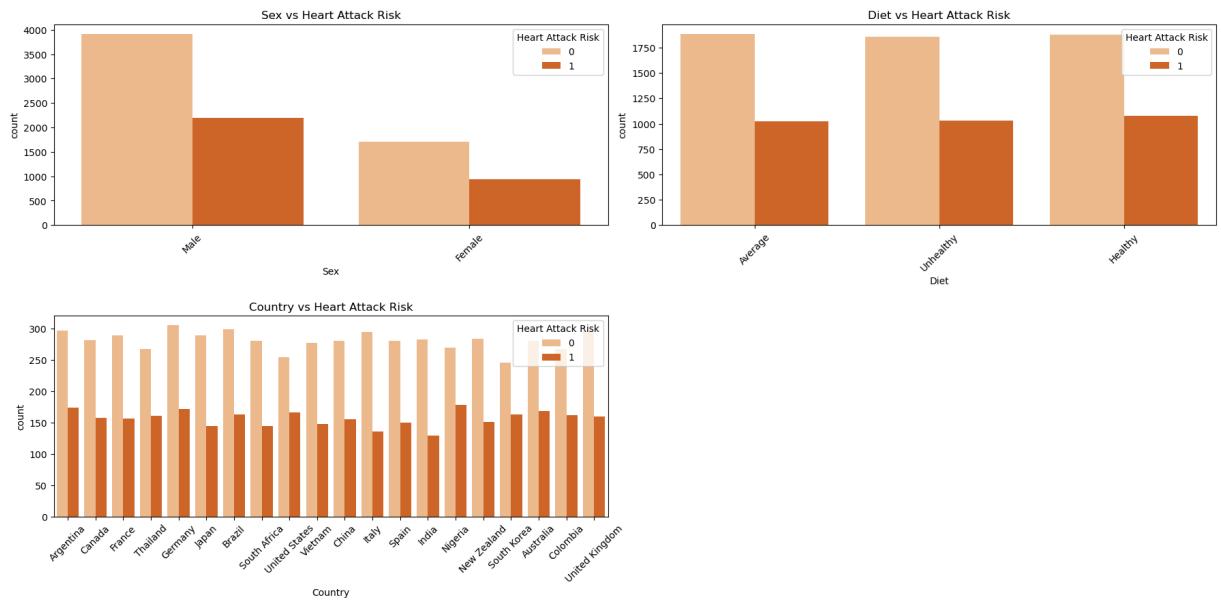


```
In [13]: # --- Correlation Heatmap --
plt.figure(figsize=(14,10))
# Select only numeric columns for correlation
corr = df.select_dtypes(include=['number']).corr()
# Focus on stronger correlations only
mask = np.triu(np.ones_like(corr, dtype=bool)) # mask upper triangle
sns.heatmap(corr, mask=mask, cmap="magma", center=0, annot=True, fmt=".2f", linewidths=cbar_kws={"shrink": 0.8})
)
plt.title("Correlation HeatMap for the Numeric Features of Heart Disease Attack", f
plt.tight_layout()
plt.show()
```



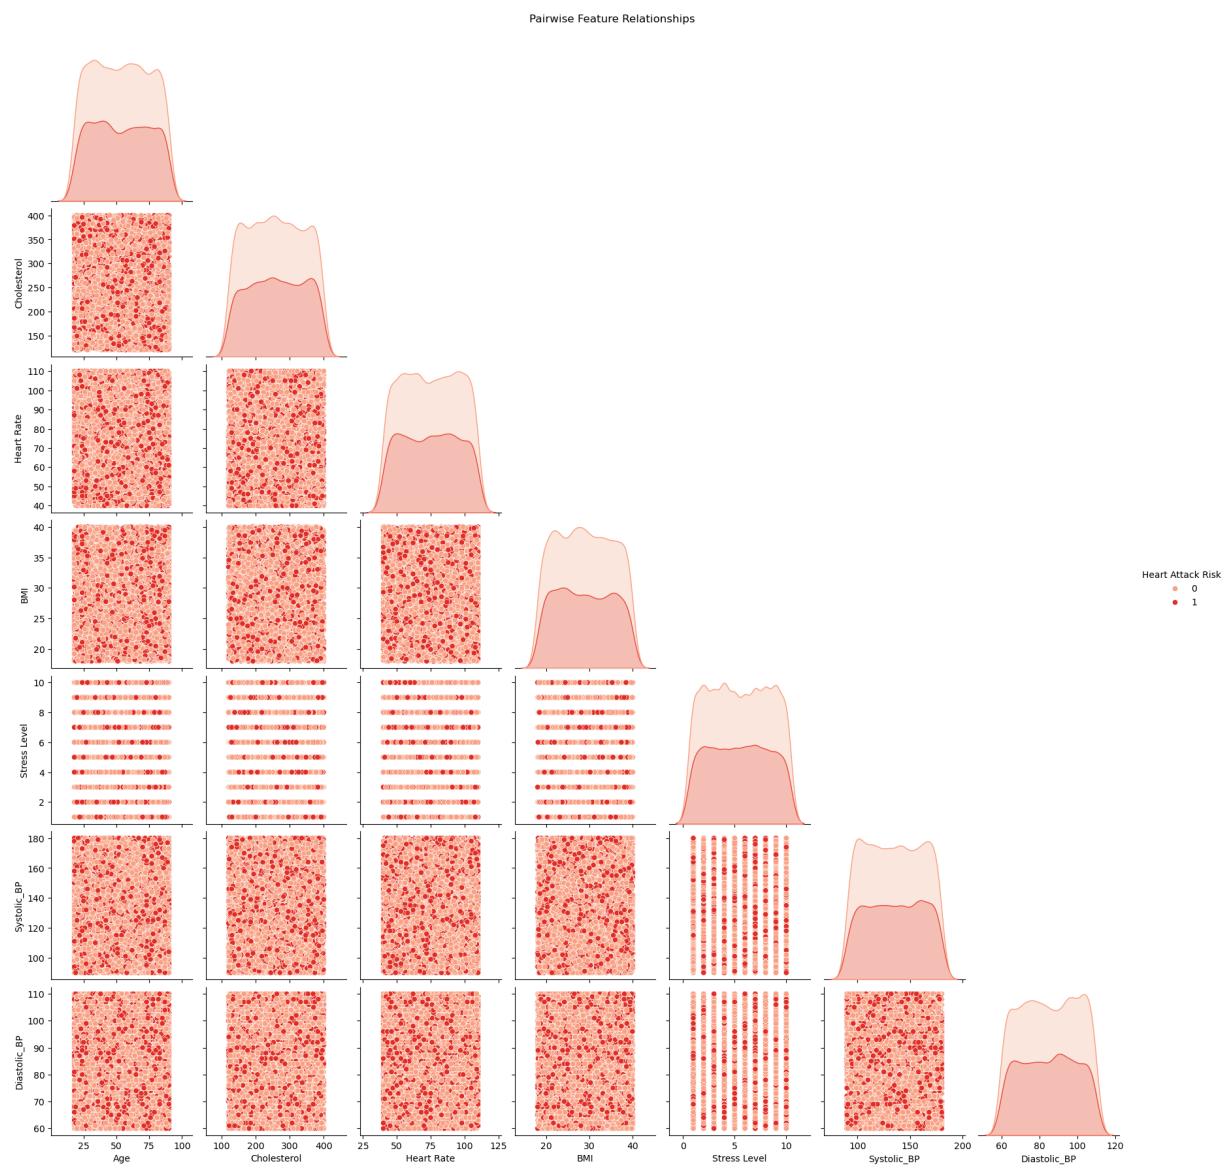
```
In [14]: # Target Distribution by Category
cat_cols = ['Sex', 'Diet', 'Country']
plt.figure(figsize=(18,12))
for i, col in enumerate(cat_cols, 1):
    plt.subplot(3, 2, i)
    sns.countplot(data=df, x=col, hue="Heart Attack Risk", palette="Oranges")
    plt.title(f"{col} vs Heart Attack Risk")
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

Heart Disease Attack

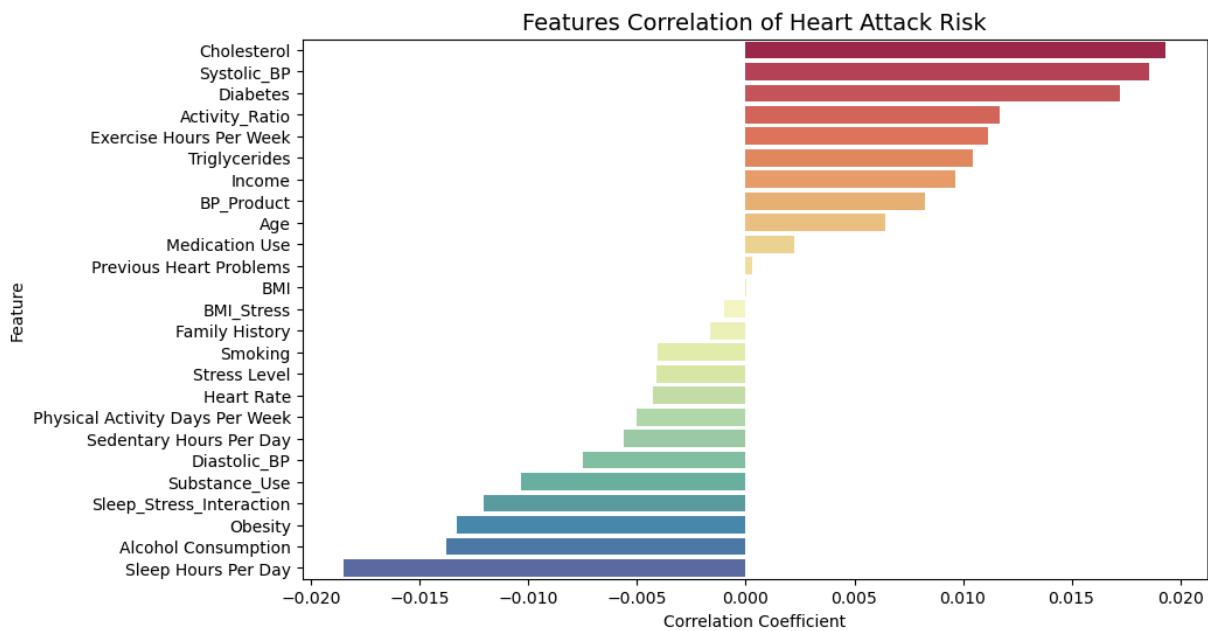


In [15]:

```
# Pairwise Relationships (Hidden Correlations)
sns.pairplot( df[['Age', 'Cholesterol', 'Heart Rate', 'BMI', 'Stress Level', 'Systolic Blood Pressure'],
hue="Heart Attack Risk", palette="Reds", diag_kind="kde", corner=True
])
plt.suptitle("Pairwise Feature Relationships", y=1.02)
plt.show()
```

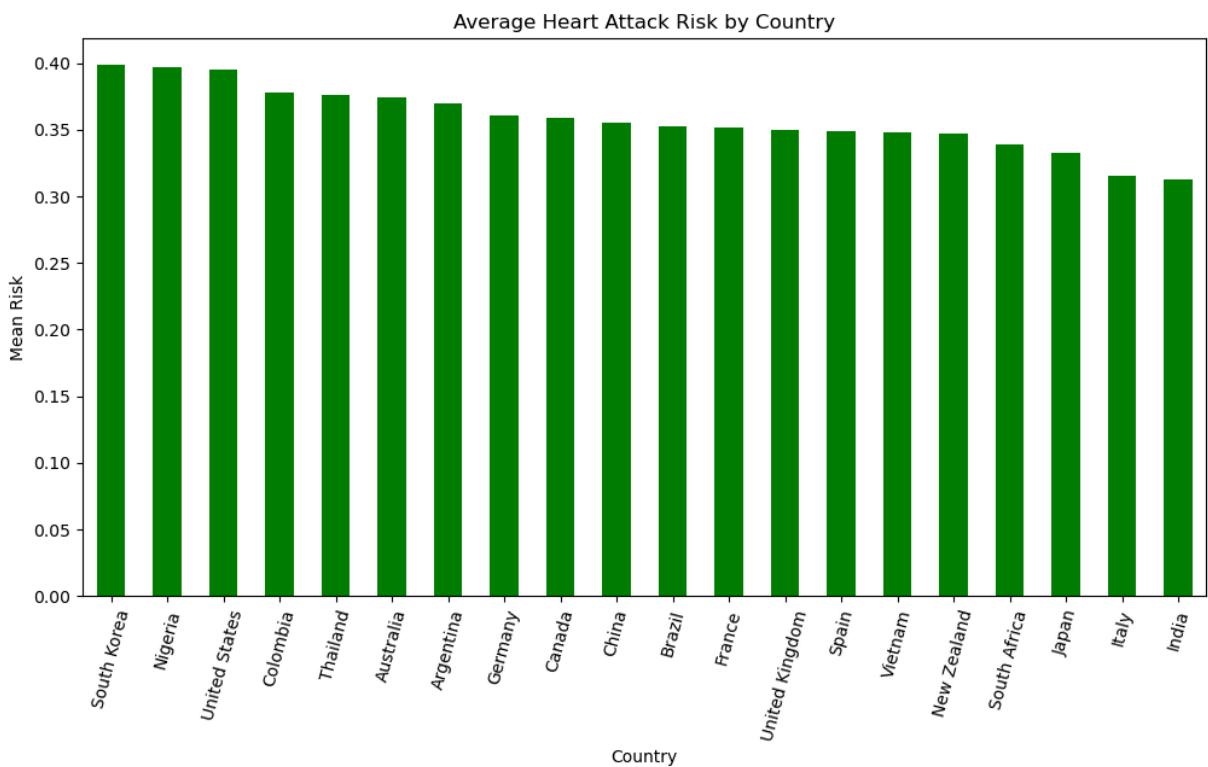


```
In [16]: # Sorted Correlation with Target
numeric_df = df.select_dtypes(include=['number'])
target_corr = (numeric_df.corr()['Heart Attack Risk']
    .drop('Heart Attack Risk')
    .sort_values(ascending=False)
)
# --- Visualize correlation strength --
plt.figure(figsize=(10, 6))
sns.barplot(x=target_corr, y=target_corr.index, palette="Spectral")
plt.title("Features Correlation of Heart Attack Risk", fontsize=14)
plt.xlabel("Correlation Coefficient")
plt.ylabel("Feature")
plt.show()
```



In [17]:

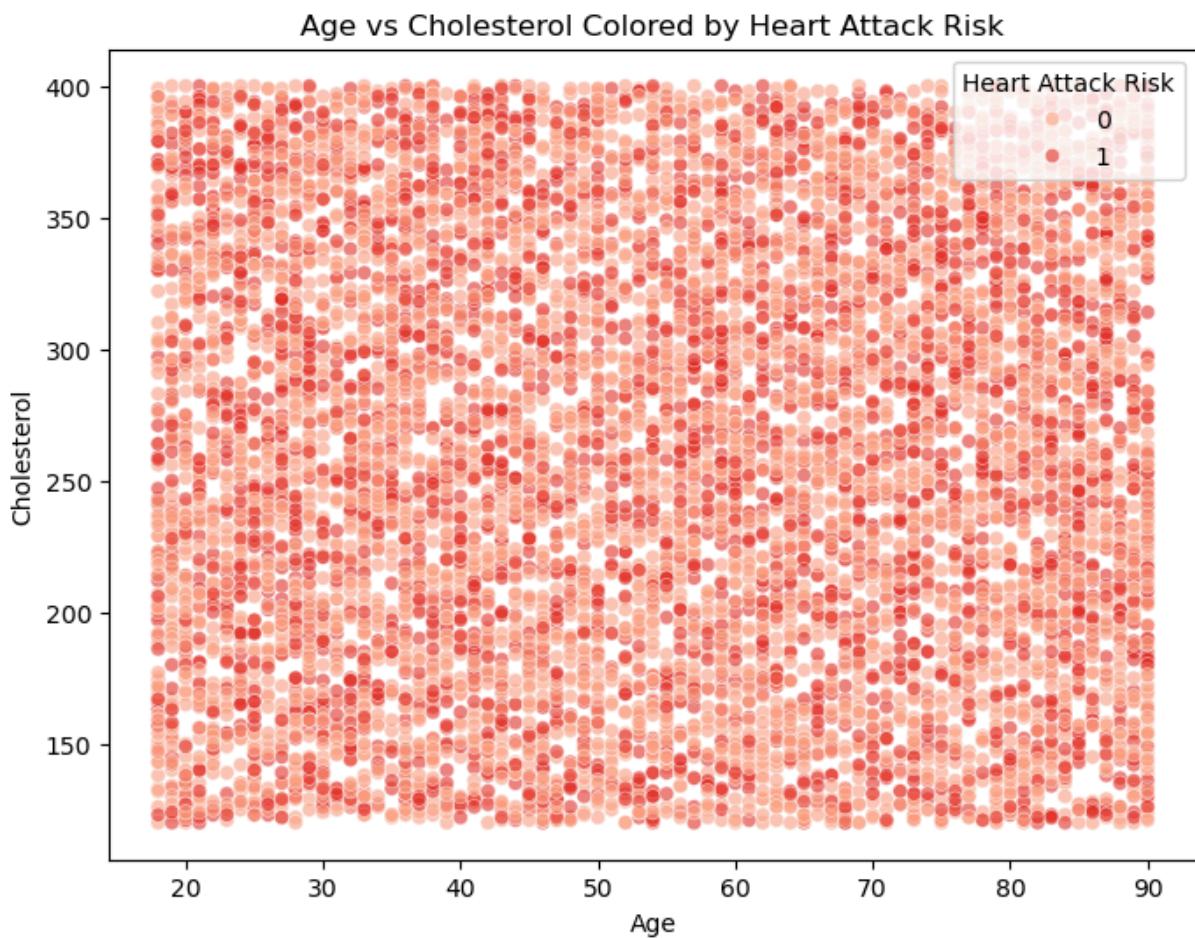
```
# Geographic Patterns(country)
risk_by_country = df.groupby('Country')['Heart Attack Risk'].mean().sort_values(as
plt.figure(figsize=(12,6))
risk_by_country.plot(kind='bar', color='green')
plt.title("Average Heart Attack Risk by Country")
plt.ylabel("Mean Risk")
plt.xticks(rotation=75)
plt.show()
```



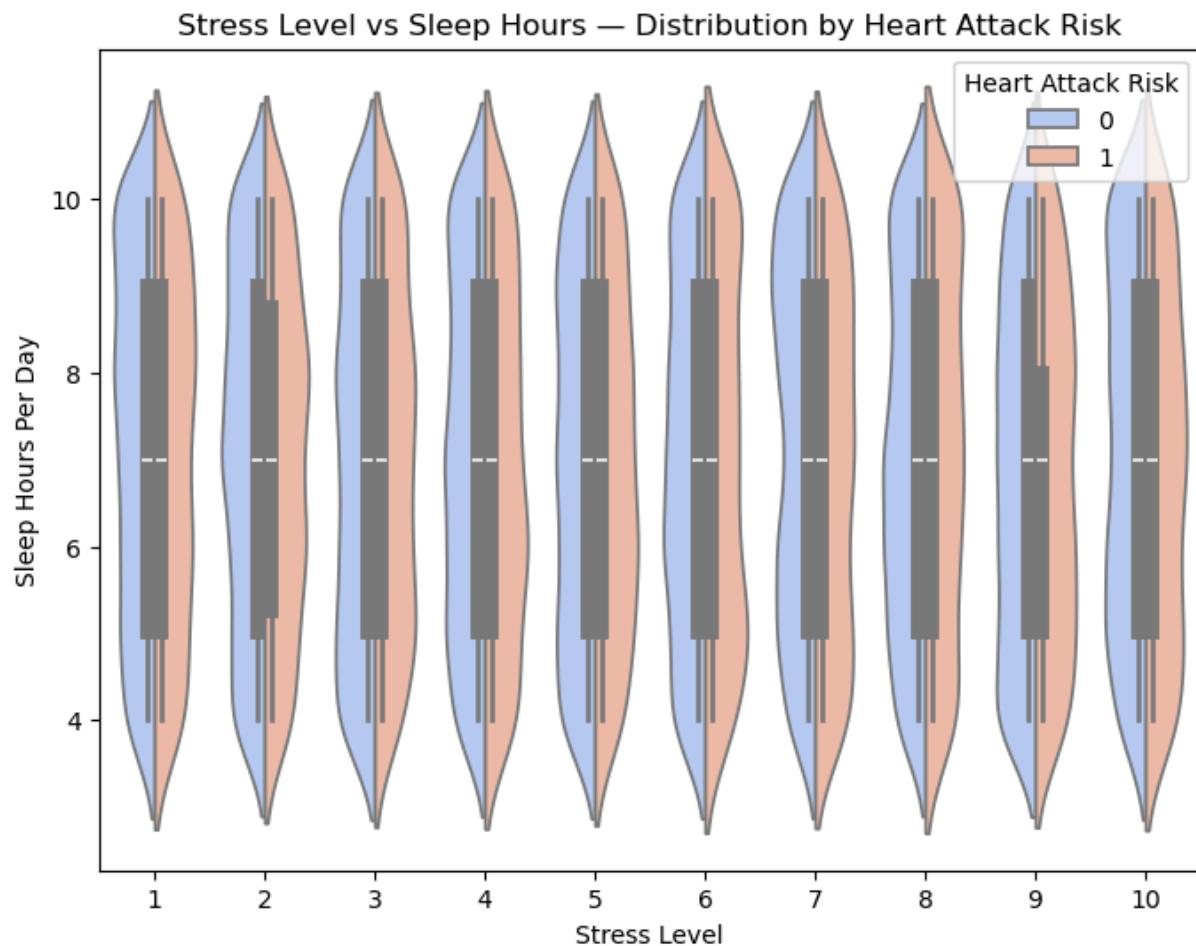
In [18]:

```
# Age vs Cholesterol Colored by Heart Attack Risk
plt.figure(figsize=(8,6))
sns.scatterplot(data= df, x='Age', y='Cholesterol', hue='Heart Attack Risk', palett
```

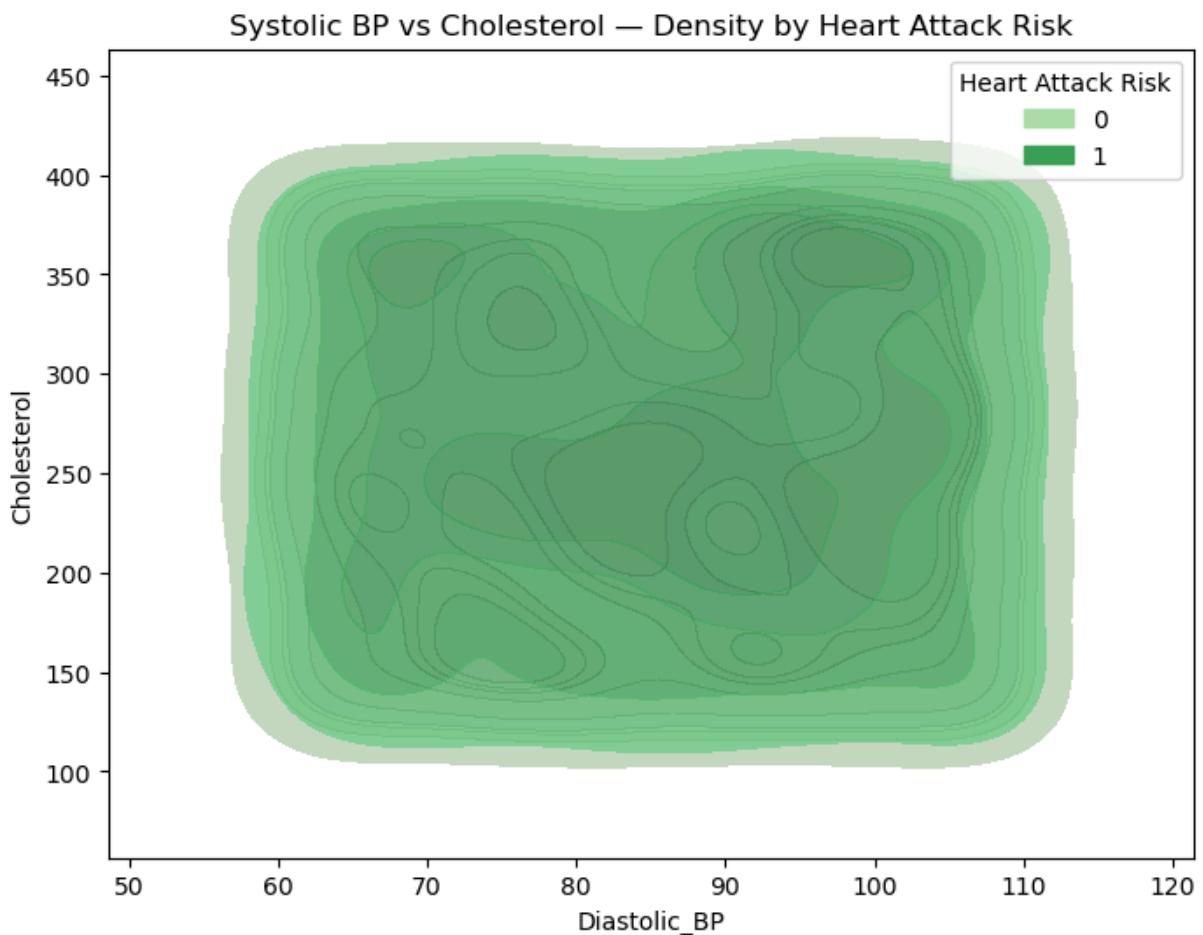
```
)  
plt.title("Age vs Cholesterol Colored by Heart Attack Risk")  
plt.show()
```



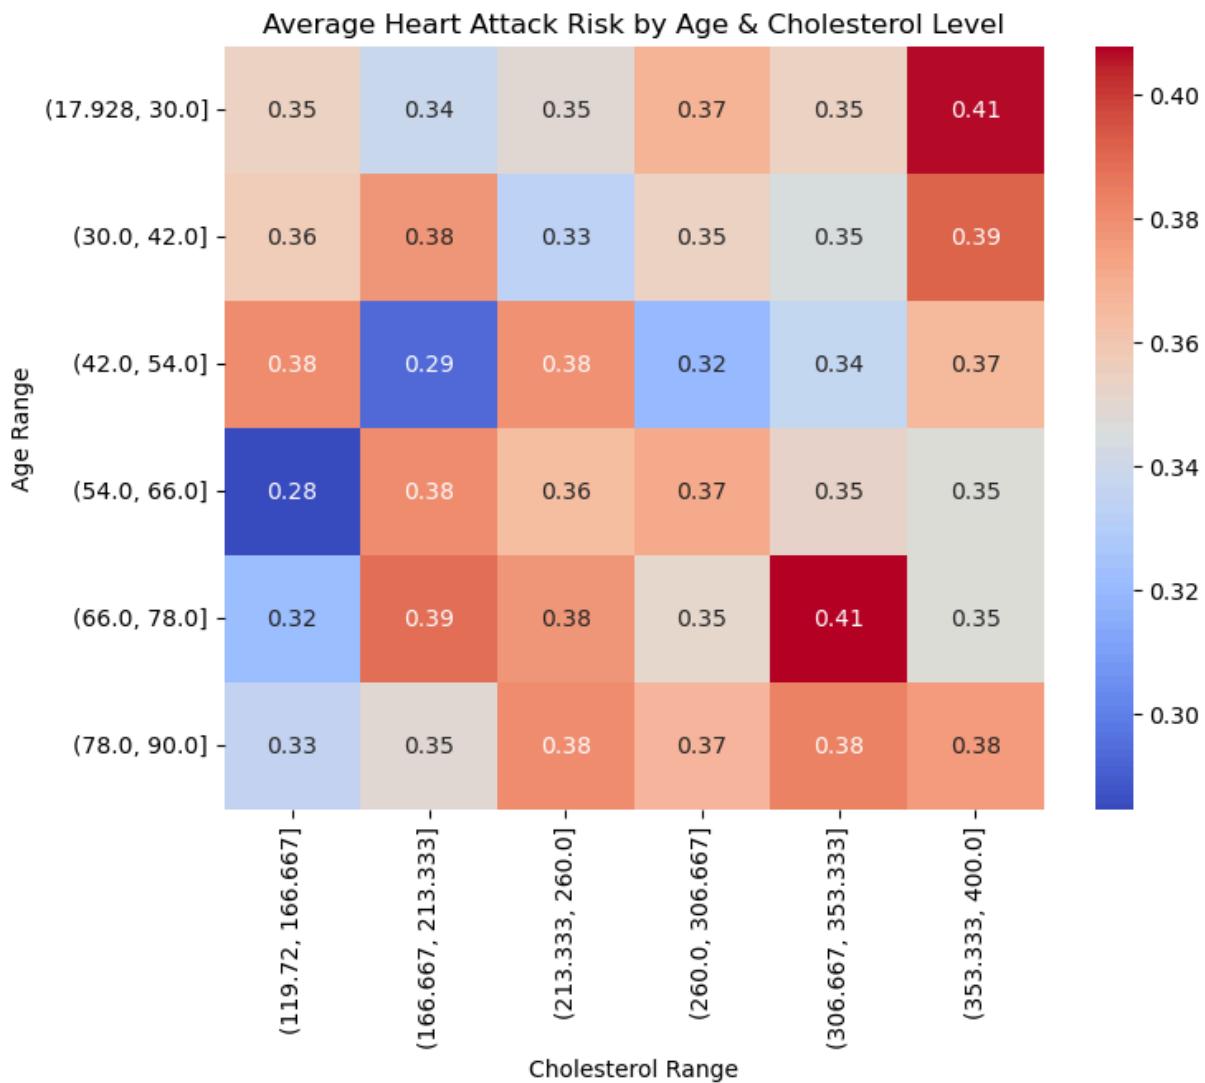
```
In [19]: # Stress Level vs Sleep Hours  
plt.figure(figsize=(8,6))  
sns.violinplot(data=df, x="Stress Level", y="Sleep Hours Per Day", hue="Heart Attack Risk")  
plt.title("Stress Level vs Sleep Hours – Distribution by Heart Attack Risk")  
plt.show()
```



```
In [20]: # Systolic BP vs Cholesterol
plt.figure(figsize=(8,6))
sns.kdeplot(data=df, x="Diastolic_BP", y="Cholesterol", hue="Heart Attack Risk", fill=True)
plt.title("Systolic BP vs Cholesterol – Density by Heart Attack Risk")
plt.show()
```

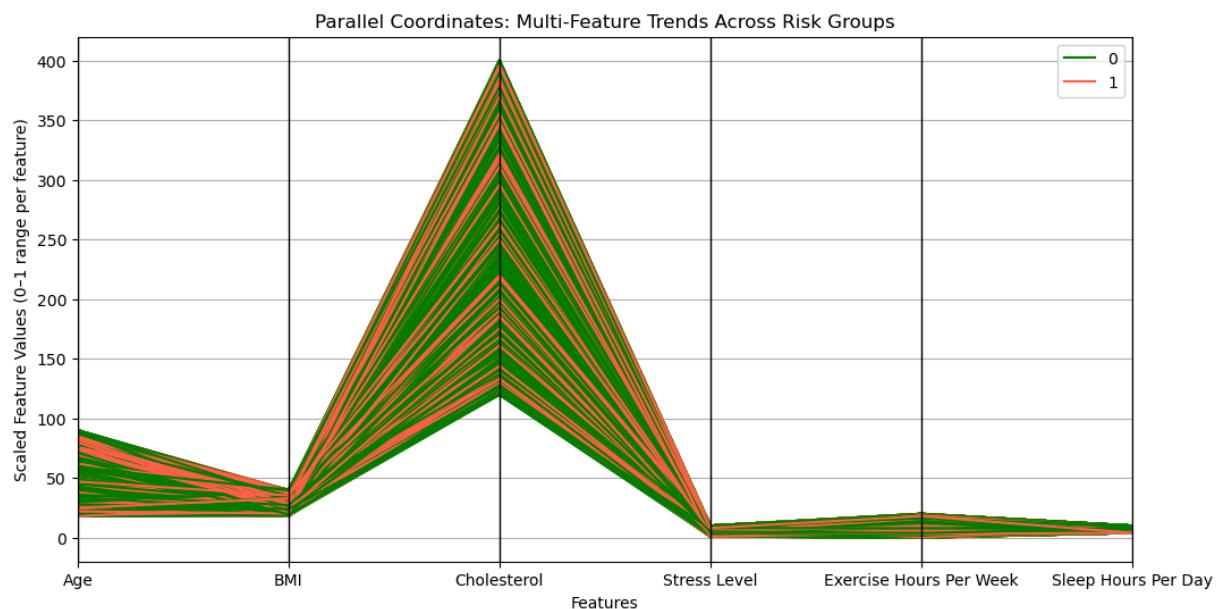


```
In [21]: # Average Heart Attack Risk by Age & Cholesterol Level
plt.figure(figsize=(8,6))
pivot = df.pivot_table(values="Heart Attack Risk",
index=pd.cut(df["Age"], bins=6),
columns=pd.cut(df["Cholesterol"], bins=6),
aggfunc="mean")
sns.heatmap(pivot, cmap="coolwarm", annot=True, fmt=".2f")
plt.title("Average Heart Attack Risk by Age & Cholesterol Level")
plt.xlabel("Cholesterol Range")
plt.ylabel("Age Range")
plt.show()
```

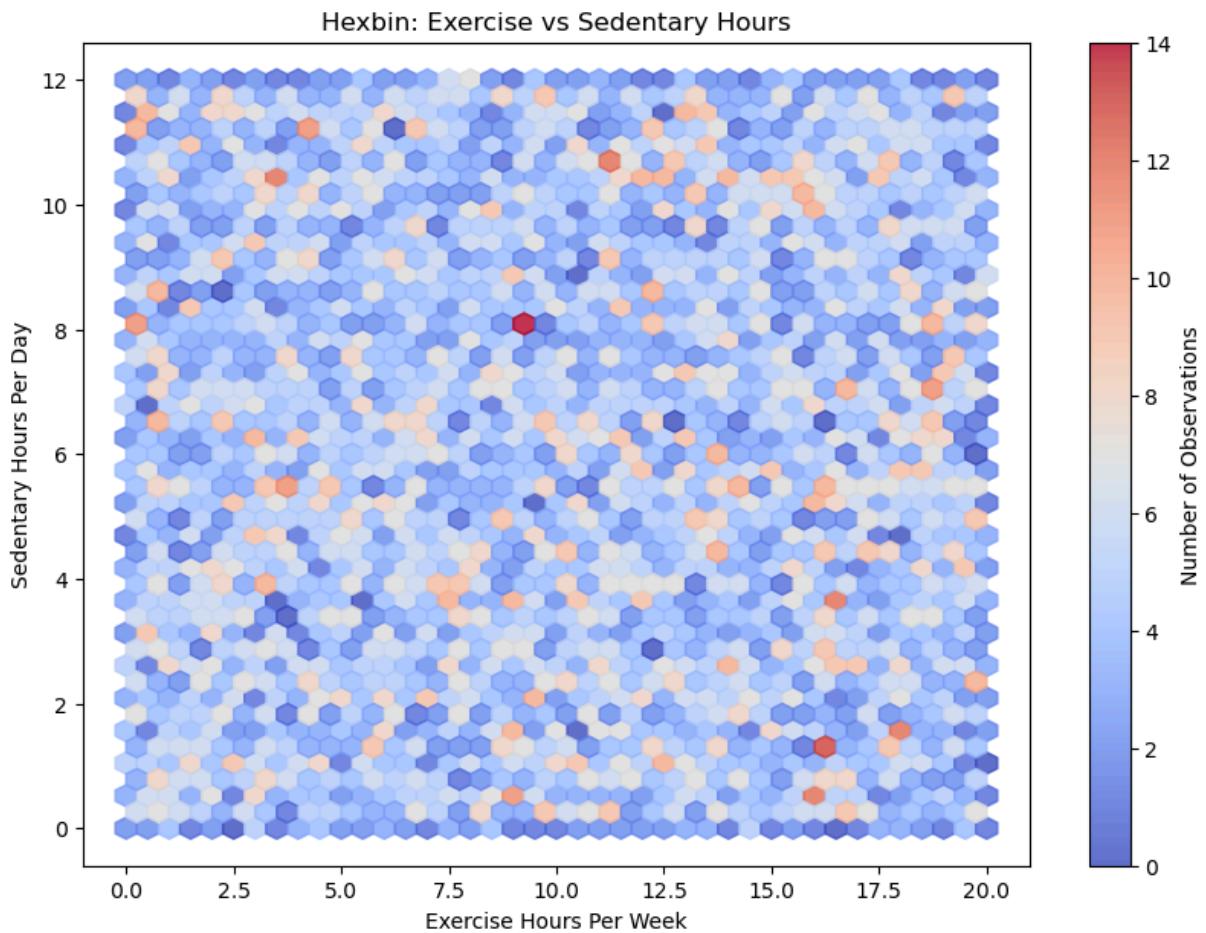


In [22]:

```
# Multi-Feature Trends Across Risk Groups
from pandas.plotting import parallel_coordinates
selected_cols = ["Age", "BMI", "Cholesterol", "Stress Level", "Exercise Hours Per
plt.figure(figsize=(12,6))
parallel_coordinates(df[selected_cols], "Heart Attack Risk", color=["green", "toma
plt.title("Parallel Coordinates: Multi-Feature Trends Across Risk Groups")
plt.ylabel("Scaled Feature Values (0-1 range per feature)")
plt.xlabel("Features")
plt.show()
```



```
In [23]: # Hexbin: Exercise vs Sedentary Hours
plt.figure(figsize=(10,7))
plt.hexbin(df["Exercise Hours Per Week"], df["Sedentary Hours Per Day"], gridsize=4)
plt.colorbar(label="Number of Observations")
plt.title("Hexbin: Exercise vs Sedentary Hours")
plt.xlabel("Exercise Hours Per Week")
plt.ylabel("Sedentary Hours Per Day")
plt.show()
```



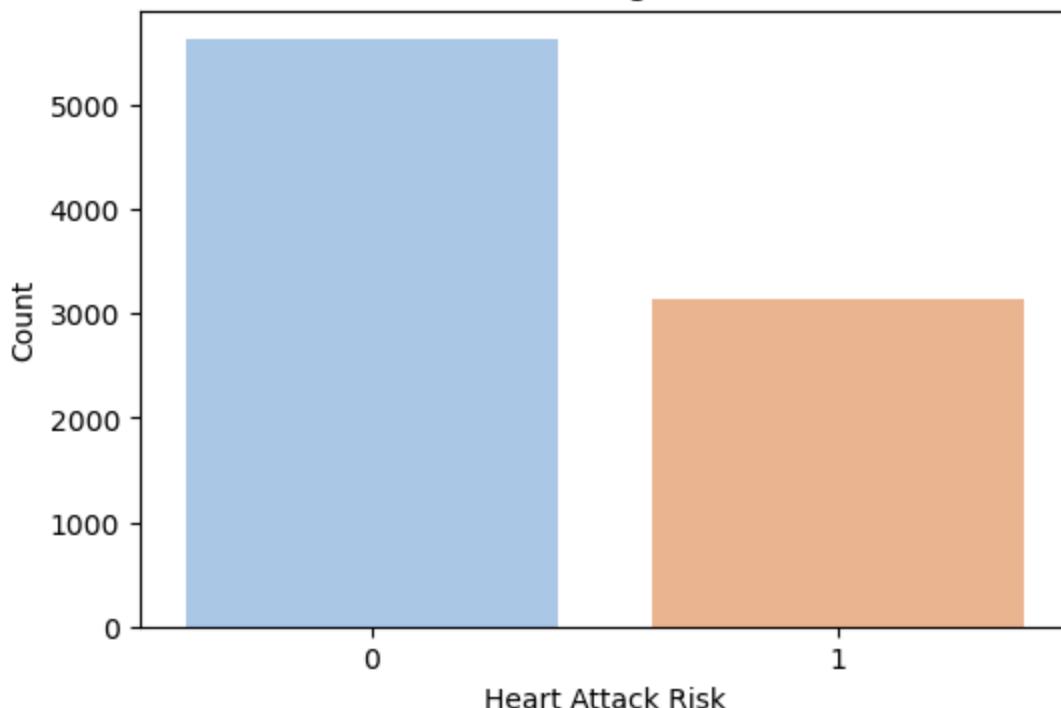
Apply SMOTE and Visualize Target Balance

Before SMOTE — check class distribution

```
In [24]: from imblearn.over_sampling import SMOTE
from collections import Counter
X = df.drop(columns=["Heart Attack Risk"])
y = df["Heart Attack Risk"]
X = pd.get_dummies(X, drop_first=True)
print("Original class distribution:")
print(Counter(y))
plt.figure(figsize=(6,4))
sns.countplot(x=y, palette="pastel")
plt.title("Before SMOTE: Target Distribution")
plt.xlabel("Heart Attack Risk")
plt.ylabel("Count")
plt.show()
```

Original class distribution:
Counter({0: 5624, 1: 3139})

Before SMOTE: Target Distribution

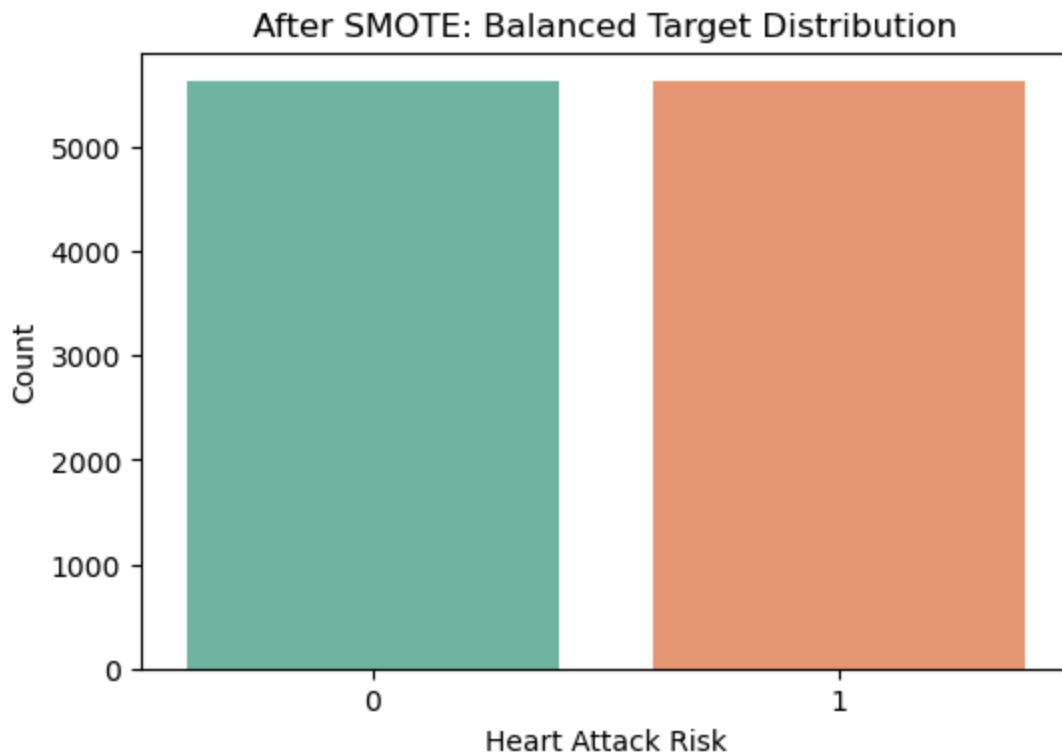


After SMOTE — check new balance

In [25]:

```
# Apply SMOTE
smote = SMOTE(random_state=42)
X_res, y_res = smote.fit_resample(X, y)
print("\n After SMOTE class distribution:")
print(Counter(y_res))
plt.figure(figsize=(6,4))
sns.countplot(x=y_res, palette="Set2")
plt.title("After SMOTE: Balanced Target Distribution")
plt.xlabel("Heart Attack Risk")
plt.ylabel("Count")
plt.show()
```

```
After SMOTE class distribution:
Counter({0: 5624, 1: 5624})
```



Create a new balanced DataFrame

```
In [26]: df_balanced = pd.concat([pd.DataFrame(X_res), pd.Series(y_res, name="Heart Attack Risk")], axis=1)
```

Out[26]:

	Age	Cholesterol	Heart Rate	Diabetes	Family History	Smoking	Obesity	Alcohol Consumption	Exercise Per Week
0	67	208	72	0	0	1	0	0	4.16
1	21	389	98	1	1	1	1	1	1.81
2	21	324	72	1	0	0	0	0	2.07
3	84	383	73	1	1	1	0	1	9.82
4	66	318	93	1	1	1	1	0	5.80
...
11243	43	201	100	1	0	1	0	0	2.20
11244	65	303	74	0	0	1	1	0	10.38
11245	25	363	62	1	0	1	1	0	9.89
11246	24	293	96	0	0	0	0	0	11.49
11247	60	312	58	1	0	1	0	1	9.96

11248 rows × 48 columns

Performing Feature Selection using Different Methods of Selection on the New Dataframe and Training Different Models with the Selected Features with Hyperparameter Tuning

importing Dependencies

In [27]:

```
# Imports
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedKFold
from sklearn.preprocessing import MinMaxScaler
from sklearn.feature_selection import SelectKBest, chi2, RFE, mutual_info_classif
from sklearn.linear_model import LogisticRegression, Lasso
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from catboost import CatBoostClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
import warnings
warnings.filterwarnings('ignore')
```

Preparing data for normalizing and traing

```
In [28]: # Split data
X = df_balanced.drop("Heart Attack Risk", axis=1)
y = df_balanced["Heart Attack Risk"]
scaler = MinMaxScaler()
X_scaled = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, str
```

Models Defination and Parameters for Tuning

```
In [29]: models_params = {
    "LogisticRegression": (LogisticRegression(max_iter=2000), {
        "C": [0.01, 0.1, 1, 10],
        "solver": ["liblinear", "lbfgs"]
    }),
    "DecisionTreeClassifier": (DecisionTreeClassifier(random_state=42), {
        "max_depth": [3, 5, 10, None],
        "min_samples_split": [2, 5, 10],
        "criterion":['gini','entropy']
    }),
    "XGBoost": (XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42), {
        "n_estimators": [100, 200],
        "learning_rate": [0.01, 0.1, 0.2],
        "max_depth": [3, 5, 7]
    }),
    "CatBoost": (CatBoostClassifier(verbose=0, random_state=42), {
        "iterations": [100, 200],
        "learning_rate": [0.01, 0.1, 0.2],
        "depth": [3, 5, 7]
    }),
    "KNN": (KNeighborsClassifier(), {
        "n_neighbors": [3, 5, 7, 9],
        "weights": ["uniform", "distance"]
    })
}
```

Defining Feature Selection Methods

```
In [30]: feature_selectors = {}
# 1. Univariate (SelectKBest)
skb = SelectKBest(score_func=chi2, k=10)
skb.fit(X_scaled, y)
```

```

feature_selectors['SelectKBest'] = X_scaled.columns[skb.get_support()]
# 2. RFE with Logistic Regression
rfe = RFE(LogisticRegression(max_iter=1000), n_features_to_select=10)
rfe.fit(X_scaled, y)
feature_selectors['RFE'] = X_scaled.columns[rfe.support_]

# 3. DecisionTree Importance
dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_scaled, y)
dt_features = X_scaled.columns[np.argsort(dt.feature_importances_)[:-10:]]
feature_selectors['RandomForest_Importance'] = dt_features

# 4. L1 Regularization (Lasso)
lasso = Lasso(alpha=0.001, max_iter=10000)
lasso.fit(X_scaled, y)
lasso_features = X_scaled.columns[np.abs(lasso.coef_) > 1e-4]
feature_selectors['L1_Lasso'] = lasso_features

# 5. Mutual Information
mi = mutual_info_classif(X_scaled, y)
mi_features = X_scaled.columns[np.argsort(mi)[-10:]]
feature_selectors['Mutual_Info'] = mi_features

```

Training and Evaluating Models

```

In [31]: results = []
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
for method, features in feature_selectors.items():
    print(f"\n Feature Selection: {method}")
    X_train_fs, X_test_fs = X_train[features], X_test[features]

    for name, (model, params) in models_params.items():
        print(f"Training {name} ...")
        grid = GridSearchCV(model, params, cv=cv, scoring='accuracy', n_jobs=-1)
        grid.fit(X_train_fs, y_train)
        best_model = grid.best_estimator_
        y_pred = best_model.predict(X_test_fs)
        acc = accuracy_score(y_test, y_pred)
        pre = precision_score(y_test, y_pred)
        recall = recall_score(y_test, y_pred)
        f1 = f1_score(y_test, y_pred)

        results.append({
            "Feature_Selection": method,
            "Model": name,
            "Best_Params": grid.best_params_,
            "Accuracy": acc,
            "Precision": pre,
            "Recall": recall,
            "F1_Score": f1
        })

```

```
Feature Selection: SelectKBest
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training KNN ...

Feature Selection: RFE
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training KNN ...

Feature Selection: RandomForest_Importance
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training KNN ...

Feature Selection: L1_Lasso
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training KNN ...

Feature Selection: Mutual_Info
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training KNN ...
```

Comparing Results

```
In [32]: results_df = pd.DataFrame(results)
results_df = results_df.sort_values(by="Accuracy", ascending=False).reset_index(drop=True)
print("\n Model Performance Summary:")
results_df
```

Model Performance Summary:

Out[32]:	Feature_Selection	Model	Best_Params	Accuracy	Precision	R
0	L1_Lasso	LogisticRegression	{'C': 10, 'solver': 'liblinear'}	0.692889	0.800554	0.51
1	L1_Lasso	XGBoost	{'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 100}	0.678222	0.714898	0.59
2	L1_Lasso	CatBoost	{'depth': 5, 'iterations': 100, 'learning_rate': 0.1}	0.665333	0.706667	0.56
3	SelectKBest	CatBoost	{'depth': 5, 'iterations': 100, 'learning_rate': 0.1}	0.640889	0.679502	0.53
4	Mutual_Info	CatBoost	{'depth': 7, 'iterations': 200, 'learning_rate': 0.1}	0.637778	0.640653	0.62
5	RandomForest_Importance	XGBoost	{'learning_rate': 0.2, 'max_depth': 7, 'n_estimators': 100}	0.634222	0.626678	0.66
6	SelectKBest	XGBoost	{'learning_rate': 0.1, 'max_depth': 5, 'n_estimators': 100}	0.631556	0.671296	0.51
7	SelectKBest	DecisionTreeClassifier	{'criterion': 'entropy', 'max_depth': 10, 'min_samples_leaf': 1}	0.628889	0.676829	0.49
8	Mutual_Info	XGBoost	{'learning_rate': 0.2, 'max_depth': 7, 'n_estimators': 100}	0.628000	0.629264	0.62
9	L1_Lasso	KNN	{'n_neighbors': 9, 'weights': 'distance'}	0.627111	0.651163	0.54
10	RandomForest_Importance	CatBoost	{'depth': 7, 'iterations': 200, 'learning_rate': 0.1}	0.625778	0.615699	0.66
11	SelectKBest	LogisticRegression	{'C': 1, 'solver': 'liblinear'}	0.622667	0.625683	0.61

	Feature_Selection	Model	Best_Params	Accuracy	Precision	R
12	RandomForest_Importance	KNN	{'n_neighbors': 3, 'weights': 'distance'}	0.621778	0.594613	0.76
13	Mutual_Info	KNN	{'n_neighbors': 9, 'weights': 'distance'}	0.601333	0.590476	0.66
14	Mutual_Info	DecisionTreeClassifier	{'criterion': 'gini', 'max_depth': 10, 'min_sa...}	0.596444	0.589007	0.63
15	SelectKBest	KNN	{'n_neighbors': 9, 'weights': 'uniform'}	0.596444	0.600743	0.57
16	L1_Lasso	DecisionTreeClassifier	{'criterion': 'entropy', 'max_depth': 5, 'min_...}	0.592889	0.616760	0.49
17	RandomForest_Importance	DecisionTreeClassifier	{'criterion': 'entropy', 'max_depth': None, 'm...}	0.582222	0.578723	0.60
18	Mutual_Info	LogisticRegression	{'C': 1, 'solver': 'lbfgs'}	0.575556	0.566719	0.64
19	RFE	CatBoost	{'depth': 3, 'iterations': 100, 'learning_rate...}	0.573778	0.564441	0.64
20	RFE	XGBoost	{'learning_rate': 0.1, 'max_depth': 3, 'n_esti...}	0.573778	0.564441	0.64
21	RFE	DecisionTreeClassifier	{'criterion': 'gini', 'max_depth': 10, 'min_sa...}	0.573778	0.564441	0.64
22	RFE	LogisticRegression	{'C': 0.01, 'solver': 'liblinear'}	0.573778	0.564441	0.64
23	RandomForest_Importance	LogisticRegression	{'C': 10, 'solver': 'lbfgs'}	0.511556	0.511168	0.52
24	RFE	KNN	{'n_neighbors': 3, 'weights': 'uniform'}	0.499111	0.499443	0.79

For Streamlit

```
In [33]: import pickle  
#save the trained model  
pickle.dump(best_model, open('heart_model.pkl', 'wb'))
```

```
In [ ]:
```